


# Clustering & Self Organizing Maps

# Outline


- Unsupervised Learning
- Clustering
  - Centroid-Based (k-means / ISODATA)
  - Distribution-Based (Gaussian Mixture Models)
  - Density-Based (DBSCAN / kNN)
  - Connectivity-Based (Hierarchical)
- Self Organizing Maps

# Unsupervised Learning


- All unlabeled data (usually because labels are expensive)
  - Learn from  $p(x)$  instead of  $p(y|x)$
  - Learn the structure of the data
- Problems:
  - Anomaly Detection
  - Density Estimation (next week)
  - Compression / Clustering



<http://blog.serverdensity.com>



[http://en.wikipedia.org/wiki/Multivariate\\_kernel\\_density\\_estimation](http://en.wikipedia.org/wiki/Multivariate_kernel_density_estimation)



<http://home.deib.polimi.it/>

# Clustering

- Goal
  - Find which instances belong to locally grouped regions (i.e. instances of distinct similarity)
- Applications
  - Marketing
  - Biology
  - Big Data
  - Image Processing
- Common Parameters
  - Distance (spread) of data
  - Number of instances in a cluster
  - Number of clusters



# Clustering Challenges

- Parameter Selection
- Scalability (Dimensionality & Cardinality)
- Feature types (nominal, ordinal, categorical)
- Arbitrarily shaped clusters
- Imbalanced Clusters


# Centroid Based Clustering

- Assume all instances belong to one and only one “centroid”
- Each centroid represents a cluster
- Centroids lie in the center of the cluster they represent
- Implies circular clusters

# k-means

- Input:  $k$ , number of clusters in data
- Procedure
  1. Initialize  $k$  cluster centroids
  2. Repeat until convergence:
    - a. Assign all instances to their nearest cluster center
    - b. Reset cluster centers to the mean of all assigned instances

# k-means (demo)






# k-means


- Expectation
  - Determine which instances belong to which centroid
- Maximization
  - Maximize the likelihood that your centroid means represent their assigned data
- Free Parameters
  - $k$  (number of centroids)
  - Distance Metric
  - Initial centroids

# k-means (when it doesn't work)



# Problems with k-means

- Number of clusters must be known
- May converge to incorrect solution (even when correct  $k$  is chosen)
- Highly dependent on initialization of centroids
- Complexity?
- Assumption that cluster center defines a cluster accurately



# ISODATA

## Iterative Self Organizing Data Analysis Technique yAy!

- Extension to k-means
  - No longer need exact number of clusters (just approximate)
- Procedure
  1. Run like k-means
  2. Split clusters whose variance is above some threshold
  3. Merge clusters that are close enough (by some threshold)
  4. Repeat
- Free Parameters:
  - $N_D$  – desired (or approximate) number of clusters
  - $N_{\text{MIN\_EX}}$  – minimum number of instances per cluster
  - $\sigma_s^2$  – maximum spread of any cluster (for splitting)
  - $D_{\text{MERGE}}$  – minimum distance between two clusters (for merging)
  - $N_{\text{MERGE}}$  – maximum number of clusters that can be merged



# ISODATA

1. Randomly initialize centroids and assign all instances to their closest centroid. Keep  $N_C$  up to date as current number of clusters
2. **Eliminate clusters** that contain less than  $N_{\text{MIN\_EX}}$  examples. Reassign examples to their nearest clusters
3. **Calculate Cluster Parameters:** For each cluster, compute the center  $\mu_k$ , the average distance,  $d_k$  between all assigned examples and  $\mu_k$ , and the cluster's axis variance  $\sigma_k^2(d^*)$  where  $d^*$  is the axis with max variance
4. **Splitting:** For all clusters with  $\sigma_k^2(d^*) > \sigma_s^2$ 
  1. If  $d_k > d_{\text{avg}}$  AND ( $N_k > 2N_{\text{MIN\_EX}}$  OR  $N_C < N_D / 2$ ) split cluster on  $d^*$ 
    1.  $\mu_{k1}(d^*) = \mu_k(d^*) + \epsilon \sigma_k^2(d^*)$        $\mu_{k2}(d^*) = \mu_k(d^*) - \epsilon \sigma_k^2(d^*)$        $0 < \epsilon < 1$
    2. Reassign examples to appropriate centers
5. **Merging:** Compute all distances between cluster centers,  $D_{ij}$ . For all,  $D_{ij} < D_{\text{MERGE}}$ 
  1. If in this iteration, clusters  $i$  and  $j$  haven't already been merged AND (not more than  $N_{\text{MERGE}}$  merges have occurred OR  $N_C > 2N_D$ )
    1. Merge clusters  $i$  and  $j$ , compute new means, assign instances accordingly, and reassign instances
6. Repeat from step 2 until convergence

# ISODATA

- Commonly used in image processing (remote sensing)
- Can eliminate clusters with few examples
- Can merge / divide clusters when needed
- Problems
  - Circular cluster assumption
  - Lots of free parameters (on which performance is highly dependent)
  - Increased computational complexity from k-means
  - Convergence is not guaranteed
- Usually ISODATA is ran multiple times with different parameters and the clustering with minimum MSE is chosen

# Distribution Based Clustering

- Assume each cluster follows a known distribution
- The parameters of the distribution are unknown and are to be estimated
- Each instance contributes to the location of all clusters




# Gaussian Mixture Models

- Assume now that each “centroid” represents a Gaussian distribution
  - We now have mean ( $\mu$ ), covariance ( $\Sigma$ ), and mixing coefficient ( $\pi$ ), for each distribution
  - Instead of assigning each instance to a center, we calculate  $p(x_i|\mu_k, \Sigma_k, \pi_k)$
- Using expectation and maximization, we maximize the total likelihood

$$p(X|\theta) = \prod_{i=1}^N p(x_i|\theta) \quad \theta = \{\mu_k, \Sigma_k, \pi_k \forall k\}$$

Why is this multiplied?





# Density Based Clustering

- Assume clusters are above some density threshold
- Connect instances in regions where the density does not fall below this threshold
- Different clusters are separated by areas of low density

# DBSCAN


- Connect instances to nearby instances if there are enough of them in a neighborhood
- **Parameters:** minimum number of instances required in a neighborhood ( $n$ ), and neighborhood radius ( $r$ )

## Procedure:

1. Initialize  $X_U$ , the set of unchecked instances, as all of the data,  $X$
2. Create a “neighborhood” centered at a test instance  $x_T \in X_U$  containing the set of instances  $X_N$  within distance  $r$  of  $x_T$ 
  - a. Remove  $x_T$  from  $X_U$
  - b. If the number of instances in  $X_N$  is larger than  $n$ , connect all  $x \in X_N$  to  $x_T$
  - c. Recursively repeat step 2 for all  $x \in X_N \cap X_U$  until all connected instances are checked
3. Repeat step 2 with a new test instance  $x_T \in X_U$

# DBSCAN

- Directly density reachable
- Density reachable
- Density connected



# DBSCAN

- Free parameters:
  - $r$  – radius of neighborhood
  - $n$  – number of instances in local region to propagate cluster
- Benefits
  - Don't need to know number of clusters
  - No assumptions on shape of clusters (what is the primary assumption?)
- Problems
  - Accidentally joining clusters that should be separate (noise) – how is this alleviated?
  - Need to approximate spread of the data
  - Complexity?




kNN?

# Hierarchical

- Types
  - **Agglomerative** (bottom-up)
  - Divisive (top-down)
- Initially consider all instances as “clusters”
- Parameter: distance measure between two clusters  $d(c_i, c_j)$ 
  - Single linkage (distance between nearest two instances)
  - Complete linkage (distance between farthest two instances)
  - Group Average (average distance between all instances)
- Group the most similar clusters, one at a time, based on  $d$

# Dendrogram



# Hierarchical Agglomerative

- How do we decide actual clustering?
- The dendrogram needs to be cut somewhere, because a bunch of different possible clusterings doesn't tell us much
- Cut criterion
  - Threshold  $d(c_i, c_j)$  or combination similarity
  - Specify number of clusters
  - $K = \underset{K'}{\operatorname{argmin}} [RSS(K') + \lambda K']$
- Problems with Hierarchical Agglomerative Clustering?




# Self Organizing Maps (SOMs)

- Teuvo Kohonen (University of Finland – 1982)
- Clustering, Visualization, & Dimensionality Reduction
- Neural Network (sort of)
  - Contains a lattice of artificial “neurons” with input weights
  - Iteratively trained like a neural network
  - No back propagation
  - No output labels
  - No activation functions
- Translate the instances in the feature space to nodes in the lattice space

# SOMs

- Connected lattice of nodes (“neurons”)
- Each node has a position
  - x, y coordinate in lattice
- Each node has a weight vector
  - Initialized randomly
  - Same dimensionality as the data
  - Map the feature space to the lattice



# SOM Procedure

## In General:

1. Randomly initialize the weight of each node
2. Sample a random instance,  $x_T$
3. Determine which node's weight is closest to  $x_T$ . This node is the best matching unit (BMU)
4. Create a kernel function,  $\theta(t)$ , centered at the BMU with radius  $\sigma_t$

$$\theta(t) = e^{-\frac{d^2}{2\sigma_t^2}}$$

$$\sigma_t = \sigma_0 e^{-\frac{t}{\lambda}}$$

5. Update all nodes within  $\sigma_t$  of the BMU

$$w_i(t+1) = w_i(t) + \theta(t)L(t)(x_T - w_i(t))$$

$$L(t) = L_0 e^{-\frac{t}{\lambda}}$$


6. Return to step 2 and repeat N times

# SOMs

- Maps data from a continuous input space to a low dimensional discrete output space
- Preserves topological structure of data
- Competitive & Cooperative
- Can be used to classify new data into the lattice
- Parameters:
  - Number of nodes and dimensionality of lattice (usually 2D)
  - Initial neighborhood radius
  - Neighborhood Kernel function
  - Time decay constant  $\lambda$
  - Initial learning rate
  - Distance measures
  - Number of iterations



# SOMs for Visualization



<http://www.cis.hut.fi/research/som-research/worldmap.html>

# References

- Clustering
  - [http://home.deib.polimi.it/matteucc/Clustering/tutorial\\_html/](http://home.deib.polimi.it/matteucc/Clustering/tutorial_html/)
  - [http://www.cs.unc.edu/~lazebnik/fall09/clustering\\_techniques\\_and\\_applications.pptx](http://www.cs.unc.edu/~lazebnik/fall09/clustering_techniques_and_applications.pptx)
  - [http://research.cs.tamu.edu/prism/lectures/pr/pr\\_115.pdf](http://research.cs.tamu.edu/prism/lectures/pr/pr_115.pdf)
  - <http://www.cs.princeton.edu/courses/archive/spr08/cos424/slides/clustering-2.pdf>
- Self Organizing Maps
  - <http://www.cs.bham.ac.uk/~jxb/NN/116.pdf>
  - <http://www.ai-junkie.com/ann/som/som1.html>
  - <http://davis.wpi.edu/~matt/courses/soms/>